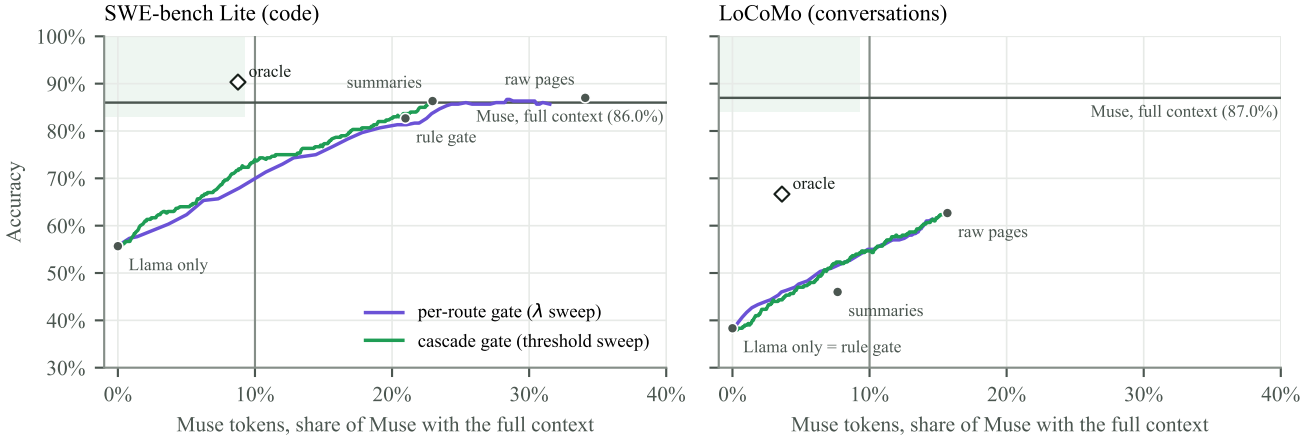
SWE-bench Lite (code)
LoCoMo (conversations)
Accuracy
100%
90%
80%
70%
60%
50%
40%
30%
oracle
summaries
raw pages
Muse, full context (86.0%)
rule gate
Llama only
per-route gate (λ sweep)
cascade gate (threshold sweep)
Muse, full context (87.0%)
oracle
raw pages
summaries
Llama only = rule gate
0%
10%
20%
30%
40%
Muse tokens, share of Muse with the full context
Muse tokens, share of Muse with the full context